

Readme

Project Category: B1AI DEVOP1

Author: Roni Fernandes Dias / François LANGE

Date: June 24, 2026

Table of Contents

Logical Databases Hosted: 5

Host Port Mapping: 5

Automated Provisioning (`database/scripts/init.sql`) 5

1. Git Commit & Pipeline Governance 8

1. Mandatory File Identification Header (CRITICAL) 8

2. Code Review & Mentorship Policy 9





3. Refactoring & Architecture Standards 9

4. Database & SQL Optimization Standards (MySQL) 9

5. Documentation Synchronization 10

6. Test-Driven Development (TDD) Policy 10

Project Antigravity (DEVOP1)

Welcome to **Project Antigravity (DEVOP1)**, an enterprise-grade Proof of Concept (PoC) demonstrating self-healing containerization, workflow automation, and unified database architectures across heterogeneous environments.

System Architecture

The following diagram illustrates the containerized development topology and data flows:

```
graph TD
    %% Define Node Styles
    classDef user fill:#e1bee7,stroke:#8e24aa,stroke-width:2px,color:#000;
    classDef git fill:#cfd8dc,stroke:#455a64,stroke-width:2px,stroke-dasharray: 4 4,color:#000;
    classDef db fill:#ffe082,stroke:#ff8f00,stroke-width:2px,color:#000;
    classDef app fill:#bbdefb,stroke:#1976d2,stroke-width:2px,color:#000;
    classDef mon fill:#b2dfdb,stroke:#00796b,stroke-width:2px,color:#000;
    classDef k8s fill:#c8e6c9,stroke:#388e3c,stroke-width:2px,color:#000;
    classDef cicd fill:#ffcc80,stroke:#e65100,stroke-width:2px,color:#000;

    User((Developer / Client)):::user
    Git[(Git Repository <br> Taiga Sync)]:::git

    subgraph DockerCompose [Docker Compose Environment]
        MySQL[(MySQL 8.0 Consolidated DB <br> Host: 3317 / Cont: 3306)]:::db

        subgraph AirflowStack [Workflow Automation]
            Airflow[Apache Airflow <br> UI & Scheduler]:::app
        end

        subgraph MonitorStack [Telemetry]
            Zabbix[Zabbix Server & UI <br> Alerts: Discord/Teams]:::mon
        end
    end

    subgraph K3s [Kubernetes K3s Cluster]
        subgraph CICD [CI/CD Pipeline]
            JenkinsMaster[Jenkins Master Pod]:::cicd
            JenkinsAgent[Dynamic Build Agents]:::cicd
        end
    end
```

```
end

subgraph Workloads [Application Deployments]
    WebApp[Web App Pods <br> Flask / React]:::k8s
    LB[Load Balancer]:::k8s
    %% Note: MetalLB and NGINX are not used currently.
end
end

%% Developer & User Workflows
User -->|API & Front-End Access| LB
LB -->|Routes Traffic| WebApp
User -->|Commits TG-XXX| Git

%% CI/CD Data Flow
Git -->|Webhooks / Polling| JenkinsMaster
JenkinsMaster -->|Provisions execution pods| JenkinsAgent
JenkinsAgent -->|Build, Push & Deploy via Helm| WebApp

%% Unified Database Connectivity
WebApp -->|Application Data <br> devop1_db| MySQL
Airflow -->|Orchestration Metadata <br> airflow_db| MySQL
Zabbix -->|Metrics & Configs <br> zabbix_db| MySQL

%% Monitoring Overlays
Zabbix -->|SNMP / Active Agents| K3s
Zabbix -->|Health Checks| MySQL
```

Technology Stack & Matrix

Layer	Technology	Purpose
User Interface	React / Vanilla HTML & CSS	Frontend for end-user interaction.
Backend API	Flask (Python 3.11)	Serves endpoints, executes `cron` / `at` system jobs.
Consolidated Database	MySQL 8.0	Single engine hosting multiple logical databases (`devop1_db` , `airflow_db`).
Workflow Automation	Apache Airflow 2.9.1	Manages execution of DAGs via `LocalExecutor`.
CI/CD Automation	Jenkins / `antigravity_bot.py`	Idempotent sync of sprints to Taiga boards and build orchestration.

Unified Database Architecture (MySQL Consolidation)

To enforce security, optimize resources, and prevent database collisions, this project **consolidates all storage services into a single MySQL 8.0 database engine**, entirely removing PostgreSQL.

Logical Databases Hosted:

1. **devop1_db**: Stores Flask backend application schemas and transactional data.
2. **airflow_db**: Serves as the Apache Airflow workflow engine's metadata repository.

Host Port Mapping:

- **Host Port:** `3317` (Exposed to avoid conflict with default local MySQL instances).
- **Container Port:** `3306` (Internal container communication is isolated).

Automated Provisioning (`database/scripts/init.sql`)

When the MySQL stack spins up for the first time, it automatically executes the initialization script:

```
CREATE DATABASE IF NOT EXISTS devop1_db;
CREATE DATABASE IF NOT EXISTS airflow_db;

CREATE USER IF NOT EXISTS 'dev_user'@'%' IDENTIFIED BY
    'your_db_password_here';
GRANT ALL PRIVILEGES ON devop1_db.* TO 'dev_user'@'%';

CREATE USER IF NOT EXISTS 'airflow'@'%' IDENTIFIED BY 'airflow';
GRANT ALL PRIVILEGES ON airflow_db.* TO 'airflow'@'%';
```

Git Workflow Standards & Cheat Sheet

We enforce strict Agile tracking and repository standards. **No changes to project files can be made without creating a task in Taiga first.** Commit messages must reference the Taiga ID (e.g., `TG-101`) to maintain clear audit logs.

Commit Format Specification

TG-<TaskID>: <Brief description in active voice>

- Detailed bullet point explaining why and what was changed
- Includes any testing verification executed

***Example:** `TG-105: Consolidate PostgreSQL to MySQL for Airflow metadata`

For all the commit a task in taiga must be associated. If the task does not exists created and add the task to a user story and a sprints.

Git Cheat Sheet

- **Create & Switch to a Feature Branch:**
`git checkout -b feature/TG-<TaskID>-description`
- **Verify Current Workspace Status:**
`git status`
- **Stage Modified & Untracked Files:**
`git add <filename>`
- **Commit with Dynamic Formatting Rules:**
`git commit -m "TG-123: Implement database consolidation"`
- **Push Branch to Origin:**
`git push -u origin feature/TG-<TaskID>-description`

Dynamic Git Filters (\$Format\$)

This repository utilizes advanced smudge/clean Git filters to dynamically inject repository metadata (author name, email, dates, commit hashes) into source code headers during checkouts, while storing them in a pristine "clean" state in Git to prevent merge conflicts.

Filter Setup Instructions

Run the setup script corresponding to your developer environment to register the filters in ``.git/config``:

- **Windows Native (PowerShell/CMD):**
`local_tools\setup_filters.bat`
- **Unix / WSL / macOS:**
`chmod +x local_tools/setup_filters.sh`
`./local_tools/setup_filters.sh`

These scripts register a **fully portable relative filter driver** that executes regardless of where your project repository is cloned:

```
git config filter.ident-dynamic.clean "python
    local_tools/git-ident-filter.py clean"
```

```
git config filter.ident-dynamic.smudge "python
  local_tools/git-ident-filter.py smudge %f"
```

Quickstart: Docker Compose Local Environment

1. **Configure the Local Environment (`.env`):**

Copy the example template file to create your local `.env`:

```
cp .env.example .env
```

Modify the `.env` parameters (e.g. database credentials, monitoring webhooks, ports) as needed.

2. **Launch the Container Stack:**

```
docker compose up -d
```

3. **Verify Service Health:**

```
docker compose ps
```

4. **Access Applications:**

- **Web Backend API:** `http://localhost:5000`
- **Apache Airflow Dashboard:** `http://localhost:8080` (Credentials: `admin`/`admin`)
- **MySQL Server:** Host `127.0.0.1`, Port `3317`, User `dev\_user` / Pass `your\_db\_password\_here`

Development Guidelines & Testing Policy

1. **Prerequisite Validation:** Every shell script and service config must programmatically verify that required environment variables exist and are not empty before proceeding with execution.
2. **Post-Task Verification:** A task cannot be resolved without a corresponding test validating its success. After executing database updates, connectivity checks must be run.
3. **WSL Integration:** WSL distributions (such as `B1AI\_DEVOP1\_LanFr144`) are linked directly to the project via `/mnt/c/...` or the home symlink (`~/devop1`), allowing easy CLI administration and verification under a Unix runtime.

Standard Repository & Skill Guidelines

This project strictly adheres to the developer standards defined across our organization's workspace. Every contribution must respect the following practices:

1. Git Commit & Pipeline Governance

- **Taiga Hook Validation:** Every commit message must start with a valid Taiga task/story ID tag (e.g., `TG-123`, `US#123`, or `[#123]`). A local git hook (`local\_tools/commit-msg`) is configured to enforce this format.
- **Pipeline Progression:** Code must progress strictly through the branches: `development` -> `test` -> `production`.
- **Segregation of Duties:** Authors promoting/merging code must not be the sole author without secondary gate review (warnings will be generated if promoting directly without review gates).

1. Mandatory File Identification Header (CRITICAL)

- **Header Tag Requirement:** Every source code, scripting, config, or text file (including ignored scratch files) must include the exact identity format at the top of the file:

```
i d e n t " @ ( # ) $ F o r m a t : { p r o j e c t _ n a m e } : { f i l
e _ n a m e } : % a n : % a e : % a d : % c n : % c e : % c d : % H : % D :
% N $ "
```

Note: In the template above, the character sequence has been intentionally formatted with spaces between each character (representing the `sed` transformation `s/./& /g`). This prevents Git's clean/smudge filters from matching, interpreting, and modifying this rule documentation file itself.

For tracked files, the Git smudge filter (`ident-dynamic`) will automatically expand the placeholder variables with real Git commit and author/commmitter data during checkouts. Untracked or ignored scratch files must still physically carry this header comment as a repository consistency requirement.

The comment syntax must match the file's language (e.g., `#` for Python/Shell/YAML/Markdown/Dockerfiles, `--` for SQL, `::` for Batch). For tracked files, the git smudge filter expands this dynamically. Ignore-listed/scratch files must carry the comment statically for structure.

To initialize a new file, place the clean version at the top of your file (legible examples are listed below in spaced-out format to prevent active smudge filter matching):

- For Python/Shell files:

```
`# i d e n t " @ ( # ) $ F o r m a t : G i t p r o j e c t n a m e : f i l e n a m e : % a n : % a e : % a d : % c n : % c e :
% c d : % H : % D : % N $ "`
```

- For SQL files:

```
`-- i d e n t " @ ( # ) $ F o r m a t : G i t p r o j e c t n a m e : f i l e n a m e : % a n : % a e : % a d : % c n : % c e :
% c d : % H : % D : % N $ "--`
```

- For Batch files:

```
:: i d e n t " @ ( # ) $ F o r m a t : G i t p r o j e c t n a m e : f i l e n a m e : % a n : % a e : % a d : % c n : % c e :
% c d : % H : % D : % N $ "::`
```

- For Markdown/YAML/Dockerfiles/XML:

```
`# i d e n t " @ ( # ) $ F o r m a t : G i t p r o j e c t n a m e : f i l e n a m e : % a n : % a e : % a d : % c n : % c e :
% c d : % H : % D : % N $ "`
```

- **Line Endings:** All files must use **LF** (Line Feed) line endings. The only exception is Windows batch scripts (*.bat), which must use **CRLF**.

- **Wiki Documentation Cadence:** Progress files under ``documentation/`` must be generated and synced to the Taiga Wiki board (``python ci-cd/antigravity_bot.py --wiki documentation/``) according to the schedule:
- Daily logs (``yyyyymmdd_daily.md``) must be compiled and synced daily.
- Plan logs (``yyyyymmdd_plan.md``) must be compiled and synced once every 2 days (excluding Sundays).
- Review logs (``yyyyymmdd_revue.md``) must be compiled and synced once every 2 days (excluding Sundays).
- **Scratch Directory Management:** Files inside ``.scratch`` are never deleted. Instead, use the archiving utility ``python local_tools/archive_scratch.py`` to move them to ``${USERPROFILE}%\keep``. If a file already exists in the destination, a 3-digit version code is appended using a semicolon (e.g., ``test_filter.py;001``, ``test_filter.py;002``).

2. Code Review & Mentorship Policy

- **Mentorship Perspective:** Senior engineers and peer reviewers adopt a constructive, instructional persona when providing code feedback.
- **Review Checklist:**
- **Correctness:** Verify the implementation solves the intended requirements.
- **Edge cases:** Handle null inputs, boundary conditions, and error states defensively.
- **Style:** Adhere strictly to project conventions.
- **Performance:** Resolve nested loops, inefficient data structures, or other potential bottlenecks.
- **Feedback Quality:** Feedback must be specific, explaining the **why** rather than just the **what**, and offering clean alternative implementations where possible.

3. Refactoring & Architecture Standards

- **DRY Principle:** Identify duplicated blocks and extract them into reusable utilities, helper functions, or libraries.
- **Complexity Reduction:** Deconstruct long, monolithic functions into small, single-purpose helper functions.
- **Micro-Files Strategy:** Prefer small, single-purpose files over monolithic structures. This ensures cleaner change tracking and simpler unit testing.
- **Behavioral Safety:** Refactored code must maintain identical external APIs and behavior without side-effects.

4. Database & SQL Optimization Standards (MySQL)

- **Performance & Locks:** Prevent row locking issues and deadlocks. Always utilize database analysis tools (like ``EXPLAIN``) for performance auditing.
- **Security & Access Control:**
- **No Hardcoded Credentials:** Application scripts must never embed plain usernames or passwords.

- **Restricted Access:** Access database objects via proxy users or restricted views owned by dedicated schemas.
- **Bind Variables:** All parameterized inputs must use bind variables. Dynamic query string concatenation is strictly forbidden.
- **Grants & Synonyms:** DDL/DML scripts must include required `GRANT` statements and `SYNONYM` configurations.
- **Transaction Management:** Auto-commit must be disabled. Explicitly control transaction scope via `COMMIT` and `ROLLBACK` blocks.
- **Syntax Standards:** Object and column names must be double-quoted (``) or back-quoted (` ` ` `) to prevent collision with reserved keywords. Avoid using keywords reserved in Oracle or standard SQL.

5. Documentation Synchronization

- **Continuous Synchronization:** Documentation, inline comments, and README instructions must be updated concurrently with any code change.
- **Architectural Impacts:** Significant system design updates must be traced and documented across the entire codebase.
- **Onboarding On-demand:** Documentation is authored assuming the reader is a new teammate, requiring maximum clarity and completeness.

6. Test-Driven Development (TDD) Policy

- **Isolation vs Integration:** Write unit tests for individual functions and integration tests for complete workflows.
- **Robust Mocks:** Mock external databases, API calls, and environment parameters using robust testing frameworks.
- **Coverage:** Strive for maximum logical coverage, specifically targeting newly modified lines of code.

9. Project Operations & Setup Manuals

- **[Deployment & Operations Manual](file:///c:/Users/your_windows_user_here/Documents/DEVOP1/antigravity/DEVOP1/docs/OPERATIONS_MANUAL.md):** The primary operational document detailing dynamic port offsetting, local setup, database backup/restores, and Kubernetes instructions.
- **[Gemini Skills Setup Guide](file:///c:/Users/your_windows_user_here/Documents/DEVOP1/antigravity/DEVOP1/docs/GEMINI_SKILLS_SETUP.md):** Explains how to place and configure Gemini agent skills to enforce development standards.
- **[Taiga Project State Export](file:///c:/Users/your_windows_user_here/Documents/DEVOP1/antigravity/DEVOP1/docs/taiga_export.json):** The JSON database snapshot showing the fully closed state of the Agile

boards, stories, and tasks.

References and Bibliography

Resource Description	URL Endpoint
![[Docker Compose]]	https://img.shields.io/badge/Orchestration-Docker_Compose-blue?logo=docker
![[MySQL 8.0]]	https://img.shields.io/badge/Database-MySQL_8.0-orange?logo=mysql&logoColor=white
![[Apache Airflow]]	https://img.shields.io/badge/Orchestrator-Apache_Airflow_2.9.1-red?logo=apache-airflow
![[K3s Kubernetes]]	https://img.shields.io/badge/Clustering-K3s_Kubernetes-green?logo=kubernetes

Page Search Index

Airflow	Page(s): 3, 4, 5, 7
Docker	Page(s): 3, 7, 8
Jenkins	Page(s): 3, 4
Kubernetes	Page(s): 3, 10
MetalLB	Page(s): 4
MySQL	Page(s): 3, 4, 5, 7, 9
RAG	Page(s): 5, 10
Telemetry	Page(s): 3
WSL	Page(s): 6, 7
Zabbix	Page(s): 3, 4